


Note : The dataset of RAVDESS and the X\_spectrograms.npy file cannot be pushed into GitHub due to large size.

Output for converting audio clips into spectrograms for training CNN:

```
0it [00:00, ?it/s]
100%|██████████| 60/60 [00:59<00:00, 1.02it/s]
100%|██████████| 60/60 [00:50<00:00, 1.20it/s]
100%|██████████| 60/60 [00:50<00:00, 1.20it/s]
100%|██████████| 60/60 [00:47<00:00, 1.27it/s]
100%|██████████| 60/60 [00:48<00:00, 1.24it/s]
100%|██████████| 60/60 [00:47<00:00, 1.25it/s]
100%|██████████| 60/60 [00:46<00:00, 1.29it/s]
100%|██████████| 60/60 [00:51<00:00, 1.16it/s]
100%|██████████| 60/60 [00:50<00:00, 1.18it/s]
100%|██████████| 60/60 [00:49<00:00, 1.21it/s]
100%|██████████| 60/60 [00:46<00:00, 1.28it/s]
100%|██████████| 60/60 [00:46<00:00, 1.30it/s]
100%|██████████| 60/60 [00:48<00:00, 1.23it/s]
100%|██████████| 60/60 [00:48<00:00, 1.23it/s]
100%|██████████| 60/60 [00:47<00:00, 1.28it/s]
100%|██████████| 60/60 [00:48<00:00, 1.25it/s]
100%|██████████| 60/60 [00:45<00:00, 1.33it/s]
100%|██████████| 60/60 [00:55<00:00, 1.08it/s]
100%|██████████| 60/60 [00:45<00:00, 1.31it/s]
100%|██████████| 60/60 [00:52<00:00, 1.15it/s]
100%|██████████| 60/60 [00:47<00:00, 1.26it/s]
100%|██████████| 60/60 [00:53<00:00, 1.13it/s]
100%|██████████| 60/60 [00:51<00:00, 1.17it/s]
100%|██████████| 60/60 [00:48<00:00, 1.23it/s]
```

Spectrogram :



## CNN Training with the obtained spectrograms :

Epoch 1,	Loss: 2.0329,	Train Acc: 19.27%,	Test Acc: 21.18%
Epoch 2,	Loss: 2.0112,	Train Acc: 18.06%,	Test Acc: 11.81%
Epoch 3,	Loss: 1.9640,	Train Acc: 24.13%,	Test Acc: 13.19%
Epoch 4,	Loss: 1.9763,	Train Acc: 22.74%,	Test Acc: 16.67%
Epoch 5,	Loss: 1.9026,	Train Acc: 24.91%,	Test Acc: 19.79%
Epoch 6,	Loss: 1.8814,	Train Acc: 28.12%,	Test Acc: 27.08%
Epoch 7,	Loss: 1.8108,	Train Acc: 31.25%,	Test Acc: 23.96%
Epoch 8,	Loss: 1.8085,	Train Acc: 31.34%,	Test Acc: 28.82%
Epoch 9,	Loss: 1.7711,	Train Acc: 30.56%,	Test Acc: 30.56%
Epoch 10,	Loss: 1.7501,	Train Acc: 34.81%,	Test Acc: 38.54%

This suggests that accuracy for testing data is low and needs further stronger models or good improvements in testing. This can be further improved.

### 1) Converting audio into text :

```
100%|██████████████████████████████████████████████████████████████████████████████| 139M/139M [00:01<00:00,  
76.5MiB/s]  
Transcribing actor: Actor_19  
Transcribing actor: Actor_21  
Transcribing actor: Actor_20  
Transcribing actor: Actor_17  
Transcribing actor: Actor_15  
Transcribing actor: Actor_18  
Transcribing actor: Actor_22
```

Transcribing actor: Actor\_23  
Transcribing actor: Actor\_24  
Transcribing actor: Actor\_16  
Transcribing actor: Actor\_06  
Transcribing actor: Actor\_13  
Transcribing actor: Actor\_10  
Transcribing actor: Actor\_07  
Transcribing actor: Actor\_14  
Transcribing actor: Actor\_11  
Transcribing actor: Actor\_05  
Transcribing actor: Actor\_08  
Transcribing actor: Actor\_09  
Transcribing actor: Actor\_12  
Transcribing actor: Actor\_04  
Transcribing actor: Actor\_01  
Transcribing actor: Actor\_03  
Transcribing actor: Actor\_02

## 2) Adding simulated texts to sentences for training text RNN :

Here in the dataset , simulated sentences play a major role in transcripts because they signify the emotion for training text RNN.

## 3) Training text RNN :

```
[nltk_data] Downloading package punkt_tab to /root/nltk_data...  
[nltk_data]   Package punkt_tab is already up-to-date!  
[nltk_data] Downloading package punkt to /root/nltk_data...  
[nltk_data]   Package punkt is already up-to-date!  
Epoch 1, Loss: 1.3968, Training Accuracy: 58.59%, Testing Accuracy: 81.25%  
Epoch 2, Loss: 0.3010, Training Accuracy: 90.36%, Testing Accuracy:  
100.00%  
Epoch 3, Loss: 0.0408, Training Accuracy: 100.00%, Testing Accuracy:  
100.00%  
Epoch 4, Loss: 0.0107, Training Accuracy: 100.00%, Testing Accuracy:  
100.00%  
Epoch 5, Loss: 0.0060, Training Accuracy: 100.00%, Testing Accuracy:  
100.00%  
Epoch 6, Loss: 0.0041, Training Accuracy: 100.00%, Testing Accuracy:  
100.00%
```

```
Epoch 7, Loss: 0.0030, Training Accuracy: 100.00%, Testing Accuracy: 100.00%
Epoch 8, Loss: 0.0023, Training Accuracy: 100.00%, Testing Accuracy: 100.00%
Epoch 9, Loss: 0.0019, Training Accuracy: 100.00%, Testing Accuracy: 100.00%
Epoch 10, Loss: 0.0016, Training Accuracy: 100.00%, Testing Accuracy: 100.00%
```

This is the output for testing data as per training data.

## Phase 2 : Early fusion

```
[nltk_data] Downloading package punkt_tab to /root/nltk_data...
[nltk_data]   Package punkt_tab is already up-to-date!
[nltk_data] Downloading package punkt to /root/nltk_data...
[nltk_data]   Package punkt is already up-to-date!
Epoch 1, Loss: 2.0587, Training Accuracy: 16.41%
Epoch 2, Loss: 1.5395, Training Accuracy: 36.28%
Epoch 3, Loss: 1.3182, Training Accuracy: 40.97%
Epoch 4, Loss: 1.1277, Training Accuracy: 47.14%
Epoch 5, Loss: 0.9037, Training Accuracy: 56.86%
Epoch 6, Loss: 0.7458, Training Accuracy: 65.10%
Epoch 7, Loss: 0.6819, Training Accuracy: 70.14%
Epoch 8, Loss: 0.5456, Training Accuracy: 75.43%
Epoch 9, Loss: 0.2755, Training Accuracy: 95.66%
Epoch 10, Loss: 0.0866, Training Accuracy: 100.00%
Test Accuracy: 100.00%
```

## Phase 2: Late Fusion

```
[nltk_data] Downloading package punkt_tab to /root/nltk_data...
[nltk_data]   Package punkt_tab is already up-to-date!
[nltk_data] Downloading package punkt to /root/nltk_data...
[nltk_data]   Package punkt is already up-to-date!
Epoch 1, Loss: 2.0724, Training Accuracy: 13.54%
Epoch 2, Loss: 1.9986, Training Accuracy: 18.14%
Epoch 3, Loss: 1.6924, Training Accuracy: 27.52%
Epoch 4, Loss: 1.6855, Training Accuracy: 26.56%
Epoch 5, Loss: 1.6543, Training Accuracy: 30.21%
Epoch 6, Loss: 1.2276, Training Accuracy: 44.97%
Epoch 7, Loss: 0.6036, Training Accuracy: 73.35%
Epoch 8, Loss: 0.2388, Training Accuracy: 87.67%
Epoch 9, Loss: 0.1579, Training Accuracy: 94.27%
Epoch 10, Loss: 0.0454, Training Accuracy: 99.39%
```

Test Accuracy: 99.65%

## Transformer of Audio Model using AST :

```
Epoch 1: Train Loss = 2.1975, Accuracy = 11.81%
Test Accuracy: 12.50%
Epoch 2: Train Loss = 2.1072, Accuracy = 13.63%
Test Accuracy: 12.50%
Epoch 3: Train Loss = 2.0848, Accuracy = 12.67%
Test Accuracy: 14.24%
Epoch 4: Train Loss = 2.0811, Accuracy = 12.15%
Test Accuracy: 12.50%
Epoch 5: Train Loss = 2.0830, Accuracy = 13.11%
Test Accuracy: 15.28%
Epoch 6: Train Loss = 2.0896, Accuracy = 12.41%
Test Accuracy: 15.28%
Epoch 7: Train Loss = 2.0914, Accuracy = 12.93%
Test Accuracy: 15.28%
Epoch 8: Train Loss = 2.0791, Accuracy = 12.85%
Test Accuracy: 14.24%
Epoch 9: Train Loss = 2.0786, Accuracy = 13.37%
Test Accuracy: 13.19%
Epoch 10: Train Loss = 2.0846, Accuracy = 14.06%
Test Accuracy: 13.54%
```

## Poetry : Dependency managing tool :

### 1) Install poetry :

```
curl -sSL https://install.python-poetry.org | python3 -
poetry --version
```

### 2) Create a new project

```
poetry new my_emotion_project
cd my_emotion_project
```

### 3) Add dependencies :

```
poetry add numpy pandas matplotlib seaborn torch  
torchvision torchaudio tensorflow transformers librosa nltk  
scikit-learn opencv-python
```

```
poetry add git+https://github.com/openai/whisper.git
```

4) Final pyproject.toml file :

Refer to file for more details

5) Install all dependencies

```
poetry install
```

6) Run the code as follows :

```
poetry run python my_emotion_project/main.py
```

Requirements.txt file :

Use “pip install -r requirements.txt” command in Ubuntu terminal to install all the required libraries for requirements.txt file.

Additional improvements in the Multi modal emotion recognition from Audio and Transcripts project :

1) Usage of a Transformer based text model instead of Text RNN like DistilBERT , BERT , Whisper encoder .

2) Many more built in features for evaluation metrics can be used instead of directly calculating accuracy .

3) Instead of using same sentences again and again for generating simulated transcripts , one thing we can use to improve semantics and vocabulary is taking some bunch of sentences for each emotion and iterate over each bunch to add them in the transcripts(original) that match with the emotion so that the RNN gets exposed to more vocabulary and it may also help for training different dataset other than the given RAVDESS.

4) We can use something like soft labels instead of directly creating an emotion map and using it to get exact emotion. (hard labels). Then we can represent each audio clip label with % of each emotion in that clip instead of just focussing on dominant emotion.

5) Using pre-trained audio models can be also good for this task.

6) Also one more thing we can use is the process of Early stopping where we stop the epochs and training if the loss is not changing much i.e converging to a point in order to save computing power and time.

7) To improve accuracy for the models, we can use deeper neural networks and also use dropout etc. We can fine-tune the transformer model.